

# RHODS Demo for Object Detection

First step deploy OCP 4.13 workshop in [demo.redhat.com](https://demo.redhat.com)

Select the Practice/ Enablement and purpose as 'conduct internal training'



## Red Hat OpenShift Container Platform 4.13 Workshop

provided by RHDP

Order



### Category

Workshops

### Product Family

Red Hat Cloud

### Provider

RHDP

### Rating

.....

### Description

Provision a cluster in Red Hat OpenShift Container Platform with basic customizations.

This asset shows how to provision an OpenShift Container Platform cluster with basic customizations for running h workshops, and similar events where no specialized environment is available.

Refer to the catalog for specialized workshops such as Containers and Cloud-Native Roadshow, OpenShift Getting Workshops. and so on.

After approx 90 mins you will get an email with login details as example below.

OpenShift Console: <https://console-openshift-console.apps.cluster-5vn9m.dynamic.opentlc.com>

OpenShift API for command line 'oc' client: <https://api.cluster-5vn9m.dynamic.opentlc.com:6443>

Download oc client from

<http://mirror.openshift.com/pub/openshift-v4/clients/ocp/4.13.1/openshift-client-linux-4.13.1.tar.gz>

Authentication via htpasswd is enabled on this cluster.

User admin with password AHRy6Bnu0ROxVUTD is cluster admin.

Normal user user1 created with password PJZteLHosOukvFn2

You can access your bastion via SSH:

ssh lab-user@ssh.ocpv01.dal10.infra.demo.redhat.com -p 30960

Enter ssh password when prompted: piscM9vplQYO

Click on the OCP Console link and use 'admin' and password given in the email but you can also see details in the demo.redhat.portal.

We will also need to ssh on to the bastion server to run some commands so be familiar with this process.

When on the OCP console install operators for RHODS and Pipelines , you should see below when completed



# Red Hat OpenShift Data Science

2.4.0 provided by Red Hat



Uninstall

## Latest version

2.4.0

## Capability level

- ☒ Basic Install
- ☐ Seamless Upgrades
- ☐ Full Lifecycle
- ☐ Deep Insights
- ☐ Auto Pilot

## Source

Red Hat

## Provider

Red Hat

## Installed Operator

This Operator has been installed on the cluster. [View it here.](#)

Red Hat OpenShift Data Science (RHODS) is a complete platform for the entire lifecycle of your AI/ML projects. It is the flagship product of OpenShift AI.

When using RHODS, your users will find all the tools they would expect from a modern AI/ML platform in an interface that is intuitive, requires no local install, and is backed by the power of your OpenShift cluster.

Your Data Scientists will feel right at home with quick and simple access to the Notebook interface they are used to. They can leverage the default Notebook Images (Including PyTorch, tensorflow, and CUDA), or add custom ones. Your MLOps engineers will be able to leverage Data Science Pipelines to easily parallelize and/or schedule the required workloads. They can then quickly serve, monitor, and update the created AI/ML models. They can do that by either using the provided out-of-the-box OpenVino Server Model Runtime or by adding their own custom serving runtime instead. These

Name Search by name...

Name	Managed Namespaces	Status	Last updated
<b>Red Hat OpenShift Pipelines</b> 1.13.0 provided by Red Hat	All Namespaces	Succeeded Up to date	6 Dec 2023, 12:50
<b>Red Hat OpenShift Data Science</b> 2.4.0 provided by Red Hat	All Namespaces	Succeeded Up to date	6 Dec 2023, 12:42

With 2.4 version you have to create the RHODS cluster, just follow the instructions and after a few minutes you will see below.

Details	YAML	Subscription	Events	<b>All instances</b>	Data Science Cluster	DSC Initialization
<b>All Instances</b>						
Filter Name Search by name...						
Name	Kind	Status	Labels	Last updated		
default-dsci	DSCInitialization	Phase  Ready	No labels	6 Dec 2023, 12:41		
rhods	DataScienceCluster	Phase  Ready	<div>app.kubernetes.io/created-by=rhods-operator</div> <div>app.kubernetes.io/instance=rhods</div> <div>app.kubernetes.io/managed-by=kustomize</div> <div>app.kubernetes.io/name=datasciencecluster</div> <div>app.kubernetes.io/part-of=rhods-operator</div>	6 Dec 2023, 12:47		

Next step will be use terminal if you have a Mac by entering the command in the email you were sent of from the [demo.redhat.com](https://demo.redhat.com), in my case I did this with command on the terminal

```
ssh lab-user@ssh.ocpv01.dal10.infra.demo.redhat.com -p 30960 ( this will be different for your setup)
```

Enter ssh password when prompted: piscM9vplQYO ( this will be different for your setup)

Then you will do first command in bold, change to the os-mlops/manifest directory after the git clone completes, then follow the next two command in bold

Note: Deploy the OpenShift Data Science operator on your OpenShift cluster. The demo pack has been tested with RHOAI 2.5.0.

For reference this is Max git repo <https://github.com/mamurak/os-mlops>

```
[lab-user@bastion ~]$ git clone https://github.com/mamurak/os-mlops
```

```
Cloning into 'os-mlops'...
```

```
remote: Enumerating objects: 3016, done.
```

```
remote: Counting objects: 100% (634/634), done.
```

```
remote: Compressing objects: 100% (267/267), done.
```

```
remote: Total 3016 (delta 433), reused 551 (delta 362), pack-reused 2382
```

```
Receiving objects: 100% (3016/3016), 15.86 MiB | 32.68 MiB/s, done.
```

```
Resolving deltas: 100% (1833/1833), done.
```

```
[lab-user@bastion ~]$ ls -lrt
```

```
total 4
```

```
drwxr-xr-x. 6 lab-user users 4096 Dec  6 07:52 os-mlops
```

```
[lab-user@bastion ~]$ cd os-mlops/
```

```
[lab-user@bastion os-mlops]$ ls
```

```
container-images  gpu-enablement.md  notebooks          odh-pachyderm-pipelines.md  README.md
```

```
rhods-kfp-modelmesh.md      rhods-workflow.png
```

```
elyra-runtime.png  manifests          odh-kfp-modelmesh.md  odh-rhods-troubleshooting.md
```

```
rhods-data-science-pipelines.md  rhods-od-workshop-instructions.md  visual-inspection-on-ocp.md
```

```
[lab-user@bastion os-mlops]$ cd manifests/
```

```
[lab-user@bastion manifests]$ oc apply -f dependencies.yaml >>Wait until the
```

```
DataScienceCluster CR has been deployed.
```

```
project.project.openshift.io/fraud-detection created
```

```
project.project.openshift.io/object-detection created
```

```
project.project.openshift.io/ray-demo created
```

```
project.project.openshift.io/spark-demo created
```

```
project.project.openshift.io/minio created
```

```
namespace/production-project created
```

```
namespace/ci-cd created
```

```
[lab-user@bastion manifests]$ oc apply -k .
```

```
role.rbac.authorization.k8s.io/spark-role created
```

```
rolebinding.rbac.authorization.k8s.io/spark-serviceaccount_name created
```

```
secret/aws-connection-fraud-detection created
```

```
secret/aws-connection-pipelines created
```

```
secret/minio-secret created
```

```
secret/aws-connection-object-detection created
```

```
secret/aws-connection-pipelines created
```

```
secret/aws-connection-pipelines created
```

```
secret/starburst-data-connection created
```

```
service/minio-service created
```

```
persistentvolumeclaim/development created
```

```
persistentvolumeclaim/fraud-detection-offline-scoring-pvc created
```

```
persistentvolumeclaim/fraud-detection-training-pvc created
```

```
persistentvolumeclaim/minio-pvc created
```

persistentvolumeclaim/development created  
persistentvolumeclaim/object-detection-training-pvc created  
persistentvolumeclaim/offline-scoring-pvc created  
persistentvolumeclaim/development created  
deployment.apps/minio created  
datasciencepipelinesapplication.datasciencepipelinesapplications.opendatahub.io/pipelines-definition created  
datasciencepipelinesapplication.datasciencepipelinesapplications.opendatahub.io/pipelines-definition created  
datasciencepipelinesapplication.datasciencepipelinesapplications.opendatahub.io/pipelines-definition created  
imagestream.image.openshift.io/code-server-cuda-data-science created  
imagestream.image.openshift.io/codeflare created  
imagestream.image.openshift.io/face-recognition created  
imagestream.image.openshift.io/fraud-detection created  
imagestream.image.openshift.io/langchain created  
imagestream.image.openshift.io/langflow created  
imagestream.image.openshift.io/monai created  
imagestream.image.openshift.io/object-detection created  
imagestream.image.openshift.io/optapy created  
imagestream.image.openshift.io/rstudio-cuda-data-science created  
imagestream.image.openshift.io/spark created  
imagestream.image.openshift.io/trino created  
imagestream.image.openshift.io/vscode-kfp created  
networkpolicy.networking.k8s.io/allow-from-all-namespaces created  
networkpolicy.networking.k8s.io/allow-from-ingress-namespace created  
networkpolicy.networking.k8s.io/allow-from-all-namespaces created  
networkpolicy.networking.k8s.io/allow-from-ingress-namespace created  
networkpolicy.networking.k8s.io/allow-from-all-namespaces created  
networkpolicy.networking.k8s.io/allow-from-ingress-namespace created  
route.route.openshift.io/minio-api created  
route.route.openshift.io/minio-ui created  
servingruntime.serving.kserve.io/ovms created  
servingruntime.serving.kserve.io/triton created  
template.template.openshift.io/mlserver created  
template.template.openshift.io/torchserve created  
template.template.openshift.io/triton created

You can leave the terminal now or leave it open as we will switch to the OCP console.

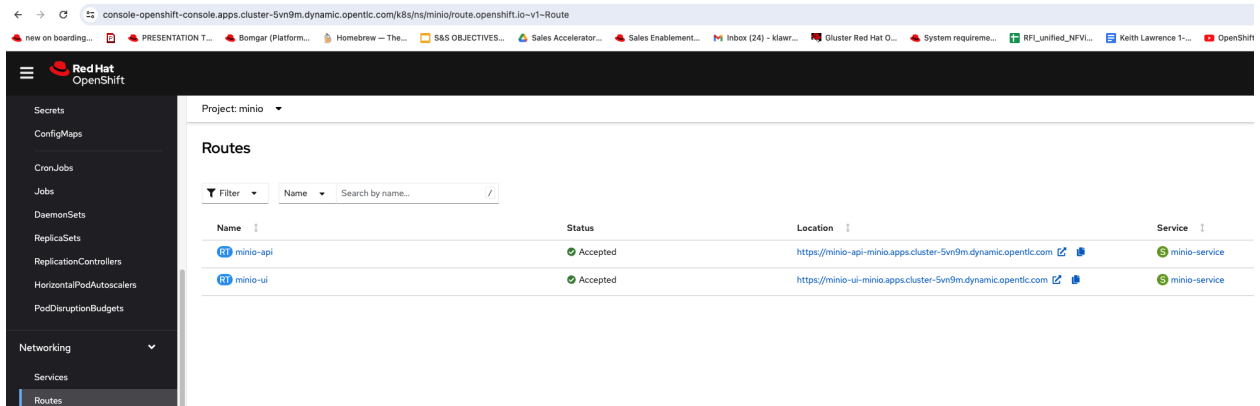
Make sure you have a tab / window open for the workshop from Codrin that will step through the object demo.

<https://codrinbucur.github.io/rhods-od-workshop/rhods-od-workshop/1-01-project-setup.html>

This is the link you will need when you are at the stage to git clone the demo but it is in the instructions above <https://github.com/mamurak/object-detection-rest.git>

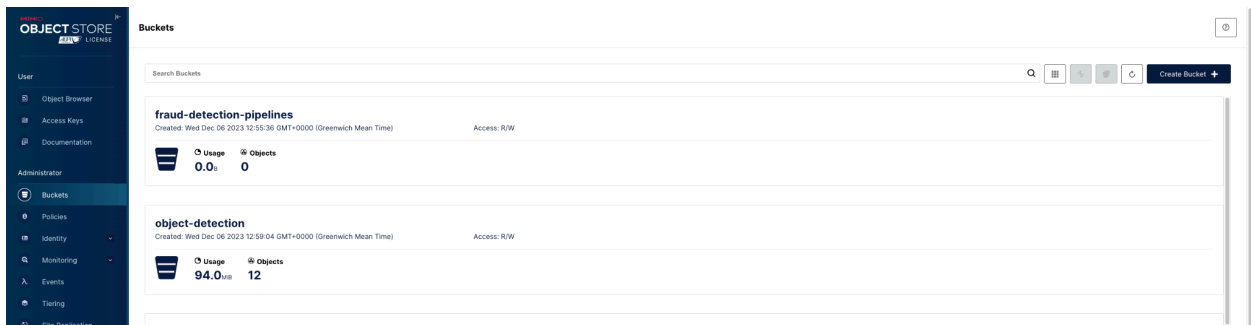
OK.....now you have executed the commands to build the Data Science projects when you were on the terminal it also set up a project for Minio which is used for S3 object storage, we need to do some manual work here..

In the OCP console we need to find the route to Minio so this is very easy and you can see from screen shot below how to achieve this, make sure you are in the Project for minio



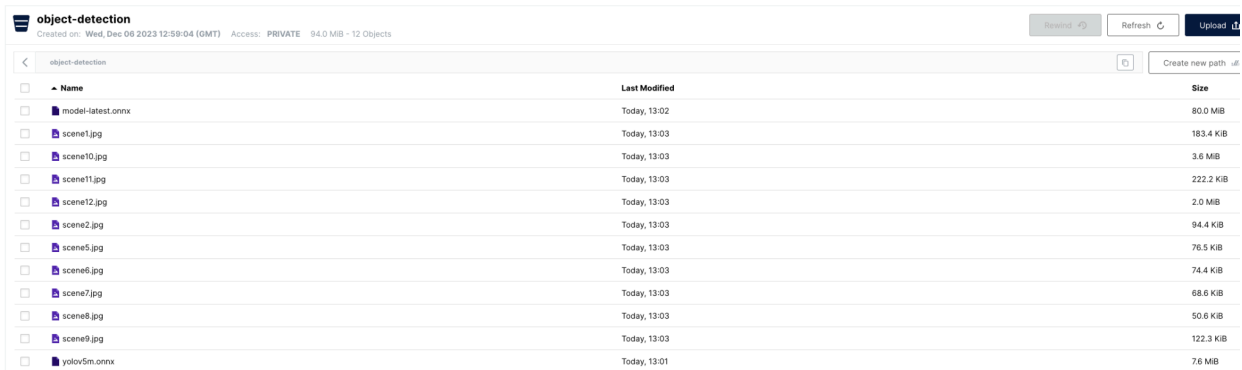
Click on the location for the minio-ui and a new tab will open, user = minio and the password = minio123

Go to the Buckets tab under Administrator, you will see some buckets but we need to create a new bucket called 'object-detection' like below



Nearly finished in this UI, we need to now go to the Object Browser under the User tab as we need to upload some files to the object-detection bucket we created.

Upload the following files I have dropped in the shared folder I will create on my google drive....<https://drive.google.com/drive/folders/1yv3HzDNsMDMV0FCrv6tNi3CkcsrLuK2c?usp=sharing>



Now we are ready to just follow the workshop instructions you have in another tab...

<https://codrinbucur.github.io/rhods-od-workshop/rhods-od-workshop/1-01-project-setup.html>